

L4Sys User Manual

Martin Unzner (munzner@os.inf.tu-dresden.de)

December 24, 2012

Abstract

This document describes how to use the L4Sys experiment suite. However, this is not a complete documentation. When in doubt, please read the source code or contact me. Still, I would like you to read this whole document before investigating further.

This is the user manual on the L4Sys generic system test framework. It provides four experiment types: GPRFlip to simulate a bit flip in a general purpose register, RATFlip to simulate a wrong association in the register allocation table, IDCFLip to corrupt a specific instruction and ALUInstrFlip to modify the behaviour of the arithmetic logic unit, so that it performs a different calculation using the same parameters.

1 Emulator Setup

These experiments work with Bochs only. This is partly due to some issues with timing — as soon as a valid model of time in the target emulator as well as an assembler/disassembler functionality in the Fail* framework are established, I would recommend a backend change, as Bochs' reliability is very limited.

To setup your system, first, you need a dedicated `bochsrc` file. It has proven handy to have a Bochs resource file or an independent Bochs instance with GUI enabled for the initial testing, however the experiments are intended to be conducted without graphical output.

2 Client Setup

All parameters of the L4Sys experiment client can be found in the file `experimentInfo.hpp`. Normally, it should not be necessary to change the program flow directly, however, if something bothers you, you are always free to take a look at `experiment.cc`, too.

2.1 Constants

Some values are constant throughout all steps of the preparation and also when the workload program is run. The most important constant is `L4SYS_BOCHS_IPS`, which has to be consistent with your `bochsrc` setting and is used for several timely calculations in the client.

2.2 Step 0: Determine the address space

First, you need to find what are the addresses of the start and end instructions of your workload program. Therefore, you should use a disassembler like *objdump* or *IDA Pro* and set `L4SYS_FUNC_ENTRY` and `L4SYS_FUNC_EXIT` accordingly. `L4SYS_NUMINSTR` is determined automatically in a later preparation step and can be ignored for now.

To test instructions that are run only in a specific address space, L4Sys uses the address space filtering mechanism of Fail*. If your fault injection campaign is not limited to a specific address space, you may set `L4SYS_ADDRESS_SPACE` to `ANY_ADDR`. Keep in mind that in that case, your instruction addresses must be unique among all those executed in the system, which in general is not the case.

To find out the value of the address space identifier, you will have to look up the value in the emulator using a debugger. In the case of Bochs, you may want to study this page, and, according to its instructions, dump the content of the CR3 control register at the beginning of your program.

2.3 Step 1: Save the initial state of the machine

Make sure `PREPARATION_STEP` is still set to 1, and you have set `L4SYS_ADDRESS_SPACE` accordingly. Now recompile and execute the framework code again, this time with the graphical user interface disabled. The experiment client runs until `L4SYS_FUNC_ENTRY` is reached and then saves the complete configuration.

2.4 Step 2: Determine the instructions to execute

For this part, it depends on how you want to conduct the injection experiments. Setting `L4SYS_FILTER_INSTRUCTIONS` stores all instructions by default, and enables the filter functionality to store only those instructions that match the filter. Each instruction in the trace requires an address plus an unsigned breakpoint counter, which means 8 bytes per instruction on a 32-bit system and 12 bytes per instruction on a 64-bit system.

If `L4SYS_FILTER_INSTRUCTIONS` is not set, the instruction to perform the fault injection at is determined by single-stepping through the program from the beginning, which is quite slow. I only recommend it for long programs, where a complete instruction trace would require several hundred megabytes of data.

No matter which method you choose, the default implementation of the campaign server reads the total instruction count from `L4SYS_NUMINSTR`. Thus, it is mandatory to set this value to the number of instructions available.

To obtain this number and optionally the instruction trace, set `PREPARATION_STEP` to 2 and recompile, then execute the experiment client. You do not have to pass parameters to Bochs any more, because the configuration is overwritten with the state saved in step 1.

After the program has finished, you will get a summary on the total of instructions executed.

If you have `L4SYS_FILTER_INSTRUCTIONS` enabled, this is not the value you look for; it merely claims how many instructions have been processed at all. To set `L4SYS_NUMINSTR` correctly, you need to look for the number before the word

accepted, which points out how many instructions have been accepted by the applied filter. Of course, if no filtering is selected, these two figures should be equal. Please contact me if that is not the case.

If `L4SYS_FILTER_INSTRUCTIONS` is disabled, you should get a statistical output on how many of the instructions were executed in userland and kernel space, respectively, but the interesting figure in this case is of course the overall sum of executed instructions.

2.5 Step 3: Determine the correct output

This is the easiest step: Set `PREPARATION_STEP` to 3, recompile the client and execute it in the target directory. It runs the complete program and logs the output. You can check the resulting file (by default `golden.out`), and if it does not comply with your expectations of a valid run, you should correct the entry and exit point, the address space or, in the worst case, your Bochs settings.

3 Campaign Setup

To setup the actual campaign, you need to edit `campaign.cc`. The full language capabilities of `AspectC++` are at your hand to define the course of your experiments; a sample covering all experiment types at random is already provided. In the experiment client, set `PREPARATION_STEP` to 0, which means there is nothing more to prepare.

After you have successfully compiled both programs, you need to start both the campaign server (`l4-sys-server`) and the experiment client. By default, they should run on the same machine, but you can adapt the `L4SysExperiment` constructor in `experiment.cc` to connect the `JobClient` to a remote server instead of `localhost`. Each experiment client processes exactly one experiment and exits. To complete your campaign, you should use the `client.sh` script in the `scripts` subdirectory of Fail*.

4 Format of the result file

When the campaign is finished, the campaign server generates a report file (by default called `l4sys.csv`) in a primitive CSV dialect. The only syntax rules are that the columns are separated by commas, that the respective data sets are separated by line breaks (`\n`), and that the cells do not contain line breaks or commas.

This section lists and describes the columns in the report generated by the campaign server, from left to right.

1. `exp_type`

Names the experiment that generated the return data. If it is none of the following, a writing error occurred:

- Unknown
- GPR Flip
- RAT Flip

- IDC Flip
- ALU Instr Flip

For *Unknown*, a debug report should be provided. If not, something went completely wrong, and you should check the logs.

2. **injection_ip**

The instruction pointer of the fault injection in lowercase hexadecimal notation. Note that the injection happens right *before* this instruction.

3. **register**

When the fault injection experiment affects a general purpose register, it is listed here. This column should have one of the following values; if it does not, a writing error occurred:

- (a) Unknown
- (b) EAX
- (c) ECX
- (d) EDX
- (e) EBX
- (f) ESP
- (g) EBP
- (h) ESI
- (i) EDI

4. **instr_offset**

The offset of the executed instruction, relative to either all executed instructions or to all listed instructions in case you applied a filter (see above). This offset includes multiple runs of the same instruction. For example, this is useful when you have loops in your program and need a rough idea how many runs your loop had executed until the injection.

5. **injection_bit**

The bit at which the injection was performed. This value is only used in GPRFlip and IDCFlip. GPRFlip inverts the bit at position **injection_bit** in the register, counted from the right. IDCFlip inverts the bit at position **injection_bit** of the current instruction, counted from the left.

6. **resulttype**

The result of the fault injection. This column should have one of the following values; if it does not, a writing error occurred:

- (a) Unknown
- (b) No effect
- (c) Incomplete execution
- (d) Crash
- (e) Silent data corruption

(f) Error

7. **resultdata**

The meaning of this field can vary for each experiment. At the moment, all of the experiments use it to store the last instruction pointer of the emulator (in decimal notation). This information can be used to determine when a fault turned into a failure.

8. **output**

The output on the EIA-232 serial line generated by the workload program. Undisplayable or reserved characters are escaped in a C conformant octal manner (e.g. \033 for the Escape character).

9. **details**

This column provides various details on the experiment run, which may help to trace errors or to reconstruct the injected fault. ALUInstrFlip uses this column to provide the opcode of the new instruction.

5 Known bugs

If you need support for more than one processor, you will have to extend the code accordingly: at the moment, when in doubt, it uses the first CPU.

6 To Be Continued

This is everything I consider important so far. If you still encounter problems you may contact me and I will try to set the record straight. Happy experimenting! :)